

## Summary

- Experienced in designing and deploying production-grade Generative AI solutions and scalable data infrastructure, delivering measurable business impact
- Proven expertise in building end-to-end RAG applications, LLM integration, and AI orchestration systems using GenAI Frameworks, and vector databases across AWS Services

## Education

- M.Tech in Software Engineering - VIT Vellore -2020

## Technical Expertise

Programming Languages: Python, SQL

Big Data Technologies: PySpark, Kafka, MongoDB, Snowflake

Cloud Technologies: Azure AI Foundry, Docker, Kubernetes, Terraform, AWS (Textract, S3, SQS)

Generative AI & ML : Generative AI, RAG, Fine-Tuning, Hugging face, Pandas, Numpy, MLFlow

Tools & Frameworks: FastAPI, Splunk, Git, Airflow, Langchain, Agno, Langgraph, Langsmith, Streamlit, Github copilot

## Industry Experience

- Banking
- Infrastructure Services | Storage

## Professional Experience

### SynkRhino: AI-Driven Data Migration Platform

- Designed an agent-based ETL migration platform that converts legacy IICS and Talend mappings into a target YAML specification, replacing manual developer-by-developer rewrites with a repeatable pipeline governed by a machine-readable contract.
- Orchestrated a multi-agent workflow — classification, SQL/expression conversion, evaluation (judge), and automated repair agent— with routing, retry and escalation controls, quality gates, and exception handling for reliable end-to-end execution.
- Built a human-in-the-loop evaluation pipeline where the judge agent validated artifacts against semantic and specification checks and a repair agent applied developer-authored rules to resolve recurring failure classes, escalating only novel failures for review.
- Built the agent pipeline with LangSmith tracing and dashboards for token/latency accounting, conversion pass rates, and judge failure categories, enabling fast failure triage across runs.

### Lead 1 Software Engineer (AI/ML Engineer) 12/2025 – 03/2026

#### Client- Experian

- Architected a GenAI application in Python and Streamlit that automates hyper-personalized marketing campaigns, cutting creative generation time.
- Built a data pipeline integrating 68k Experian behavioral metrics into Azure AI Search to ground AI content in verifiable data, and applied Azure AI Foundry across the workflow for guardrails and hallucination checks.
- Led the architectural refactoring of the Arabica evaluation framework, using Python Protocols and dependency injection to decouple the core orchestration loop and support any tech stack across teams.
- Delivered a configuration-driven design letting internal teams pass model endpoints, custom prompts, and other variables via config files without touching framework code.

## **Lead 1 Software Engineer (AI/ML Engineer) 05/2025-12/2025 Client – Global Payments**

### **Multi-Agent RAG ChatBot**

- Designed a modular Agentic workflow decoupling the application into specialised workers (language detection, guardrails, response agents), improving system stability and debuggability versus monolithic chains.
- Implemented model routing across the Agentic workflow, directing complex tasks to higher-capability LLMs and simpler tasks to lower-cost models — balancing response quality against cost and latency.
- Built a cost-efficient RAG ingestion pipeline with an intelligent document router that selectively triggers AWS Textract for complex forms, achieving a 90% reduction in document-processing cost.
- Designed the prompt layer governing agent behavior at each conversation stage — intent classification, language detection, and response generation — with LangSmith tracing and guardrail checks scoring retrieval relevance and response quality to catch incomplete or hallucinated output before it reached the user

### **.Gen AI Engineer - NetApp Jan 2024 - April 2025**

- Designed and implemented RAG based applications that increased the productivity of customer support team by 50% using Azure OpenAI-4.
- Streamlined incident management by developing a Splunk-integrated LLM system that autonomously analyzes error logs and drafts support tickets with resolution or steps, significantly reducing mean time to resolution (MTTR).
- Built scalable data lake infrastructure supporting AI/ML workflows, processing multi-TB daily volumes with Airflow and PySpark to power predictive analytics and GenAI applications.

### **Data engineer II NetApp May2023 - Jan 2024**

- Engineered and optimized multi-cloud data pipelines leveraging Apache Spark, ensuring transformed customer data meets stringent business criteria and seamless cross-platform data quality.
- Worked on comprehensive data pipeline integrating diverse data sources like API's, different files from cloud and leveraging technologies such as Spark, Monogdb, Kafka, snowflake.

### **Data engineer I NetApp June 2020 - May 2023**

- Worked on Service Management Suite Health Monitoring System to enhance ticket tracking and monitoring, enabling improved service delivery and issue resolution.
- Managed and scaled a Data Lake, ingesting and preprocessing data from various partner sources to provide curated datasets for the NetApp DataScience teams, facilitating their analytical and modeling effort

### **Engineering Intern | NetApp Oct 2019 - June 2020**

- Worked on Modernizing a hyper-converged infrastructure management platform by migrating microservices from Docker Swarm to Kubernetes managed services
- Automated Kubernetes cluster provisioning for development and CI/CD pipelines using Terraform, improving infrastructure scalability and deployment efficiency.